

Introduction to Machine Learning

Chap II: Regression



Linear regression

- Basic concepts
- Non linear data and basis functions

Model optimization

- How to control your model
- How to optimize your model


Model minimization

- Gradient descent

Bonus: Likelihood and regression (*)

- Variable/feature: x , weight: w
- Vector: variables $\mathbf{x} = \begin{pmatrix} x_1 \\ \vdots \\ x_n \end{pmatrix}$; weights $\mathbf{w} = \begin{pmatrix} w_1 \\ \vdots \\ w_n \end{pmatrix}$
- Vector (transpose): $\mathbf{x}^T = (x_1, \dots, x_n)$
- Dot product : $\mathbf{w} \cdot \mathbf{x} = \mathbf{w}^T \mathbf{x} = \sum_{i=1}^n w_i x_i = w_1 x_1 + w_2 x_2 + \dots w_n x_n.$
- Sequence of p vectors: $\{\mathbf{x}_j\}_{(j=1..p)}$
- Matrix (size $n \times m$): $\mathbf{M} = \begin{pmatrix} a_{11} & \cdots & a_{1m} \\ \vdots & \ddots & \vdots \\ a_{n1} & \cdots & a_{nm} \end{pmatrix}$

Linear regression



I DON'T TRUST LINEAR REGRESSIONS WHEN IT'S HARDER
TO GUESS THE DIRECTION OF THE CORRELATION FROM THE
SCATTER PLOT THAN TO FIND NEW CONSTELLATIONS ON IT.


Simple case: 1 dimensional data

Training dataset

- N observations of **feature** $\mathbf{x} = \{x_1, \dots, x_N\}$
- N **Target** values $\mathbf{t} = \{t_1, \dots, t_N\}$

Prediction model: straight line

$$y(x, \mathbf{w}) = y(x; w_0, w_1) = w_0 + w_1 x$$



Weights determined by minimizing an **Error function E**

- also called **Cost function** or **Loss function** (i.e what is the consequence of your error !)

Common choice: sum of **square distance** between function and target:

$$E(w_0, w_1) = \sum_{i=1}^N \{y(x_i; w_0, w_1) - t_i\}^2$$


Simple case: 1 dimensional data

Training dataset

- N observations of **feature** $\mathbf{x} = \{x_1, \dots, x_N\}$
- N **Target** values $\mathbf{t} = \{t_1, \dots, t_N\}$

Prediction model: straight line

$$y(x, \mathbf{w}) = y(x; w_0, w_1) = w_0 + w_1 x$$



Here **optimal weights** can be calculated **analytically** (not always possible !)

$$E(w_0, w_1) = \sum_{i=1}^N \{y(x_i; w_0, w_1) - t_i\}^2$$
$$\begin{cases} \frac{\partial E(w_0, w_1)}{\partial w_0} = 0 \\ \frac{\partial E(w_0, w_1)}{\partial w_1} = 0 \end{cases} \Leftrightarrow \begin{cases} w_1 = \frac{\text{cov}(x, t)}{\text{var}(x)} = r \frac{\sigma(t)}{\sigma(x)} \\ w_0 = \bar{t} - r \frac{\sigma(t)}{\sigma(x)} \bar{x} \end{cases}$$

(r: correlation factor between x and t)


Simple case: 1 dimensional data

Training and testing

- **Training:** use dataset to determine **weights** w_0 and w_1
- **Testing:** check compatibility of $y(x, \mathbf{w})$ on a new dataset

Measure of **compatibility**: root mean squared error (RMS)

$$E_{RMS} = \sqrt{\frac{1}{N} \sum_{i=1}^N \{y(x_i, \mathbf{w}) - t_i\}^2} = \sqrt{\frac{E(\mathbf{w})}{N}}$$



Generalization: multidimensional data

Dataset (p x 1 data)

- **N** observations of **p**-dimensions **features**

$$\{\mathbf{x}_i\}_{i=1..N} = \{\mathbb{R}^p\} = \left\{ \begin{pmatrix} x_1 \\ \vdots \\ x_p \end{pmatrix} \right\}$$


- **N target** values $\mathbf{t} = \{t_1, \dots, t_N\}$

Fit function: multidimensional plane

- Linear function with $p+1$ weights: \mathbf{w}

$$y(\mathbf{x}, \mathbf{w}) = w_0 + \mathbf{w}^T \mathbf{x} = w_0 + w_1 x_1 + w_2 x_2 + \dots w_p x_p.$$

bias term



Generalization: multidimensional data


Dataset (p x q data)

- **N** observations of **p**-dimensions **features**

$$\{\mathbf{x}_i\}_{i=1..N} = \{\mathbb{R}^p\} = \left\{ \begin{pmatrix} x_1 \\ \vdots \\ x_p \end{pmatrix} \right\}$$

- **N target** values of **q**-dimensions


$$\{\mathbf{t}_i\}_{i=1..N} = \{\mathbb{R}^q\} = \left\{ \begin{pmatrix} t_1 \\ \vdots \\ t_q \end{pmatrix} \right\}$$



Fit functions:

$$\begin{pmatrix} y_1(\mathbf{x}, \mathbf{w}) \\ \vdots \\ y_q(\mathbf{x}, \mathbf{w}) \end{pmatrix} = \underbrace{\begin{pmatrix} w_{01} \\ \vdots \\ w_{0q} \end{pmatrix}}_{\text{bias terms}} + \begin{pmatrix} w_{11} & \cdots & w_{1p} \\ \vdots & \ddots & \vdots \\ w_{q1} & \cdots & w_{qp} \end{pmatrix} \begin{pmatrix} x_1 \\ \vdots \\ x_p \end{pmatrix}$$

What if your data is not linear ?



→ use basis functions

Linear basis function models

Apply **M non-linear basis functions** ϕ to input feature \mathbf{x} :

$$\mathbf{x} \longrightarrow \begin{pmatrix} \phi_1(\mathbf{x}) \\ \vdots \\ \phi_M(\mathbf{x}) \end{pmatrix} \quad \phi_j(\mathbf{x}): \text{basis function}$$

The regression function $y(\mathbf{x}, \mathbf{w})$ then become non-linear function of \mathbf{x} :

$$y(\mathbf{x}, \mathbf{w}) = w_0 + \sum_{i=1}^M w_i \phi_i(\mathbf{x}) = w_0 + w_1 \phi_1(\mathbf{x}) + \cdots + w_M \phi_M(\mathbf{x})$$

These functions are called **linear models** because they are linear in \mathbf{w} .

For high number of dimensions linear models suffer from **limitations**, and other approaches (as Neural Networks) are more suited.


Linear basis function models

Polynomial basis functions (1D)

$$\phi_j(x) = x^j$$

$$y(x, \mathbf{w}) = \sum_{j=0}^{M-1} w_j x^j$$

Global functions of input variable
→ a small change in x affects all
basis functions



Linear basis function models

Gaussian basis functions (1D)

$$\phi_j(x) = e^{-\frac{(x-\mu_j)^2}{2\sigma^2}}$$


$$y(x, \mathbf{w}) = \sum_{j=0}^{M-1} w_j e^{-\frac{(x-\mu_j)^2}{2\sigma^2}}$$

Parameters:

μ_j (location) and σ (width)

Normalization is not relevant.

local functions of input variable
→ a small change in x mostly
affects nearby basis functions



Linear basis function models

Sigmoidal basis functions (1D)

$$\phi_j(x) = \sigma \left(\frac{(x - \mu_j)}{s} \right)$$


with

$$\sigma(a) = \frac{1}{1 + \exp(-a)}$$

Parameters:

μ_j (location) and s (slope)


local functions of input variable
→ a small change in x mostly
affects nearby basis functions



Example: polynomial curve fitting

Training dataset

- N observations of $\mathbf{x} = \{x_1, \dots, x_N\}$: uniformly spaced in $[0,1]$
- Target values $\mathbf{t} = \{t_1, \dots, t_N\}$: $\sin(2\pi x) + \text{Gaussian noise}$



Dummy example but could be e.g. temperature (t) evolution over 1 day (x)

Polynomial curve fitting

Fit function

- Polynomial function of degree M , with coefficients $\mathbf{w} = (w_1, \dots, w_M)^T$

$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$

- Non-linear function of x , but linear function of $\mathbf{w} \rightarrow$ **linear model**
- Values of coefficient obtained by **minimizing** an **error function**
- Sum of the square of the errors $E(\mathbf{w})$**


$$E(\mathbf{w}) = \sum_{i=1}^N \{y(x_i, \mathbf{w}) - t_i\}^2$$

Minimization



Fitted weights \mathbf{w}^*

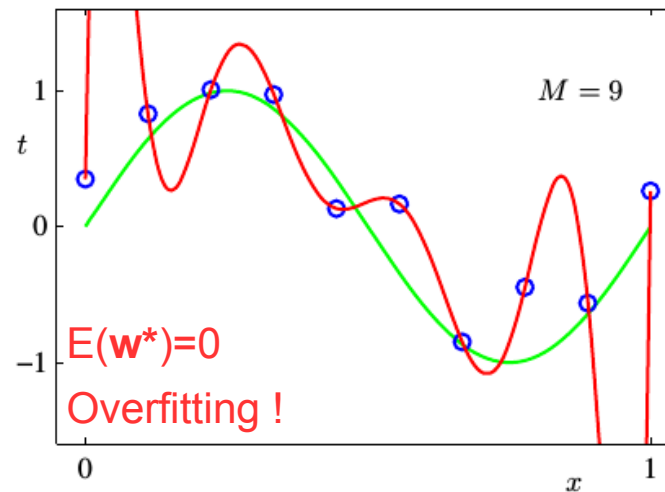



$E(\mathbf{w}^*)$





Overfitting

$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \cdots + w_Mx^M = \sum_{j=0}^M w_jx^j$$




It is instructive to look at the **fitted weights** for various cases: when M increases the coefficient become **fine tuned** to data by developing large positive and negative values.


	$M = 0$	$M = 1$	$M = 3$	$M = 9$
w_0^*	0.19	0.82	0.31	0.35
w_1^*		-1.27	7.99	232.37
w_2^*			-25.43	-5321.83
w_3^*			17.37	48568.31
w_4^*				-231639.30
w_5^*				640042.26
w_6^*				-1061800.52
w_7^*				1042400.18
w_8^*				-557682.99
w_9^*				125201.43

Root mean squared error (RMS)

$$E_{RMS} = \sqrt{\frac{1}{N} \sum_{i=1}^N \{y(x_i, \mathbf{w}) - t_i\}^2} = \sqrt{\frac{E(\mathbf{w})}{N}}$$





What if your model completely fails ?



→ try to regularize your model

Overfitting really depends on **N** data and **M** parameters.



How can we constrain the fitted parameter into reasonable values ?

→ **Regularization** techniques can be a solution.

Add **penalization term** to error function in order to **constrain** parameters \mathbf{w} .

→ Simple penalization: **ridge regression** (L2 norm)

Constrains weight to be not too large .

$$\tilde{E}(\mathbf{w}) = \sum_{i=1}^N \{y(x_i, \mathbf{w}) - t_i\}^2 + \lambda \|\mathbf{w}\|^2$$

where $\|\mathbf{w}\|^2 = \mathbf{w}^T \mathbf{w} = w_0^2 + \dots + w_M^2$

and λ : parameter that governs the importance of regularization

Other choices



- **Lasso** regression (L1 norm): $\|\mathbf{w}\| = |w_0| + \dots + |w_M|$
Reduce number of weights (set some of them to 0)
- **Elastic net**: L1 + L2 norm

General regularization term is of the form:


$$\tilde{E}(\mathbf{w}) = \sum_{i=1}^N \{y(x_i, \mathbf{w}) - t_i\}^2 + \lambda \sum_{j=1}^M |w_j|^q$$

Minimizing this error function is equivalent to minimizing the unregularized sum-of-square error with the constraint


$$\sum_{j=1}^M |w_j|^q \leq \eta$$




Lasso
(L1 norm)



Ridge
(L2 norm)





Regularization (*)



The optimum value for the parameter vector w is denoted by w^* .
The lasso gives a sparse solution in which $w_1^* = 0$.

Regularization




	$\ln \lambda = -\infty$	$\ln \lambda = -18$	$\ln \lambda = 0$
w_0^*	0.35	0.35	0.13
w_1^*	232.37	4.74	-0.05
w_2^*	-5321.83	-0.77	-0.06
w_3^*	48568.31	-31.97	-0.05
w_4^*	-231639.30	-3.89	-0.03
w_5^*	640042.26	55.28	-0.02
w_6^*	-1061800.52	41.32	-0.01
w_7^*	1042400.18	-45.95	-0.00
w_8^*	-557682.99	-91.53	0.00
w_9^*	125201.43	72.68	0.01

Effect of L2 norm regularization

- $\ln \lambda = -\infty$: no regularization
- $\ln \lambda = -18$: suppressed overfitting
- $\ln \lambda = 0$: fit too constrained

Regularization

L2 norm: $\lambda ||\mathbf{w}||^2$

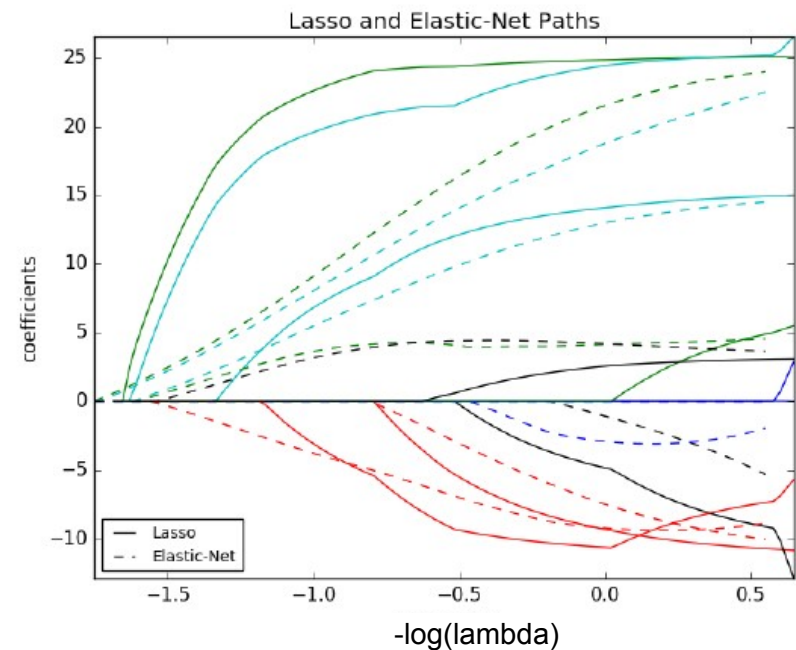


←

More constraints Less constraints

Affects value of coefficients
(shrinkage)

L1 norm: $\lambda ||\mathbf{w}||$



←

More constraints Less constraints


Affects number of coefficients
(sparsity)



K-fold cross-validation

Divide data in K groups, use K-1 for training and test on left-over group

Rinse and repeat K times



Cross-validation error: $CV = \frac{1}{K} \sum E_i$


Choose the set of hyper-parameters (ex λ) that give the smallest CV.

Drawback: can be very **time consuming** ...

K-fold cross-validation


Divide data in K groups, use $K-1$ for training and test on left-over group

Rinse and repeat K times



Cross-validation curve

What is a good model ?




[figure: kdnuggets.com]

The Bias-Variance decomposition

Training dataset

- N observations of **feature** $\mathbf{x} = \{x_1, \dots, x_N\}$
- N **Target** values $\mathbf{t} = \{t_1, \dots, t_N\}$

We assume that \mathbf{t} are distributed following a function: $t_i = f(x_i) + \epsilon$



Noise
(Mean 0,
variance σ^2)

→ We want to find $y(x)$ that approximates true function $f(x)$

The Bias-Variance decomposition (*)

As before we determine $y(x)$ by **minimizing**: $\sum_{i=1}^N \{y(x_i, \mathbf{w}) - t_i\}^2$
over the **training** dataset

The **expected error** for a **new test sample \mathbf{x}** can be decomposed as:

- Data **noise**: **minimal error** of the model
- **Bias** in the model: error caused by **model assumptions**
- **Variance** of model: how much $y(x)$ depends on **structure of data**

The Bias-Variance decomposition (*)

As before we determine $y(x)$ by **minimizing**: $\sum_{i=1}^N \{y(x_i, \mathbf{w}) - t_i\}^2$
over the **training** dataset

The **expected error** for a **new test sample x** can be decomposed as:


- Data **noise**: **minimal error** of the model
- **Bias** in the model: error caused by **model assumptions**
- **Variance** of model: how much $y(x)$ depends on **structure of data**

Mean squared error :

$$\mathbb{E}[(y(x) - f(x))^2] = \boxed{\sigma^2} + \boxed{(\mathbb{E}[y(x)] - f(x))^2} + \boxed{\mathbb{E}[(y(x) - \mathbb{E}[y(x)])^2]}$$

The expectation \mathbb{E} is calculated over
different choices of the training dataset

The Bias-Variance decomposition



Simple models **under-fit**: deviate from data (high bias) but not influenced by structure of data (low variance)


Complex models **over-fit**: small deviation from data (low bias) but very sensitive to data fluctuations (high variance)





How can we **minimize** the error function for complex cases (ex: when there is no analytic solution) ?

→ **Solution: Gradient descent**

Iteratively move in the direction of steepest descent as defined by the negative of the gradient of the error function



Descend along the error function to find a (local) minimum:



Direction of descent:

→ (negative of the **gradient** of the error function) × (**learning rate**)

Example: fit N data points with linear function: $y(x, \mathbf{w}) = w_0 + w_1 x$

Error function and its derivatives

$$E(w_0, w_1) = \sum_{i=1}^N \{y(x_i, \mathbf{w}) - t_i\}^2 = \sum_{i=1}^N \{(w_0 + w_1 x_i) - t_i\}^2$$

$$\longrightarrow \begin{cases} \frac{\partial E(w_0, w_1)}{\partial w_0} = \sum_{i=1}^N 2 \{(w_0 + w_1 x_i) - t_i\} \\ \frac{\partial E(w_0, w_1)}{\partial w_1} = \sum_{i=1}^N 2 x_i \{(w_0 + w_1 x_i) - t_i\} \end{cases}$$

Example: fit N data points with linear function: $y(x, \mathbf{w}) = w_0 + w_1 x$

Error function and its derivatives

$$E(w_0, w_1) = \sum_{i=1}^N \{y(x_i, \mathbf{w}) - t_i\}^2 = \sum_{i=1}^N \{(w_0 + w_1 x_i) - t_i\}^2$$

$$\longrightarrow \begin{cases} \frac{\partial E(w_0, w_1)}{\partial w_0} = \sum_{i=1}^N 2 \{(w_0 + w_1 x_i) - t_i\} \\ \frac{\partial E(w_0, w_1)}{\partial w_1} = \sum_{i=1}^N 2 x_i \{(w_0 + w_1 x_i) - t_i\} \end{cases}$$

Iterative update rule:

$$w_0^{(k)} \rightarrow w_0^{(k+1)} = w_0^{(k)} - \frac{\partial E(w_0, w_1)}{\partial w_0} \times \eta$$

$$w_1^{(k)} \rightarrow w_1^{(k+1)} = w_1^{(k)} - \frac{\partial E(w_0, w_1)}{\partial w_1} \times \eta$$


k: iteration number

η : learning rate


Repeat until convergence

Gradient descent

Input data: $\{x_i, t_i\}$




Error function




$$E(w_0, w_1) = \sum_{i=1}^N \{(w_0 + w_1 x_i) - t_i\}^2$$

Gradient descent

Input data: $\{x_i, t_i\}$



Gradient descent



$$w_0^{(k)} \rightarrow w_0^{(k+1)} = w_0^{(k)} - \frac{\partial E(w_0, w_1)}{\partial w_0} \times \eta$$

$$w_1^{(k)} \rightarrow w_1^{(k+1)} = w_1^{(k)} - \frac{\partial E(w_0, w_1)}{\partial w_1} \times \eta$$

1000 iterations


learning rate $\eta = 0.05$

Stochastic gradient descent


Gradient descent can be **computationally costly** for large N since the gradient is calculated over full training set.

→ **Solution: Stochastic gradient descent**


Compute gradient on a small **batch** of events (can be 1 event):

$$\begin{cases} \frac{\partial E(w_0, w_1)}{\partial w_0} = \sum_{i \in N} 2 \{ (w_0 + w_1 x_i) - t_i \} \\ \frac{\partial E(w_0, w_1)}{\partial w_1} = \sum_{i \in N} 2 x_i \{ (w_0 + w_1 x_i) - t_i \} \end{cases}$$


Stochastic gradient descent



Gradient calculated on **1**
(random) event at each step



Gradient calculated on **10**
(random) events at each step



Likelihood and regression (*)


Likelihood

Consider \mathbf{N} measurements of x distributed along a given probability law $p(x)$.

$$\mathbf{x} = (x_1, \dots, x_N)$$

where values x_i are **independent and identically distributed** (i.i.d).

Ex: Normal (a.k.a Gaussian) law with 2 parameters: mean μ and variance σ^2



$$\mathcal{N}(x|\mu, \sigma^2) = \frac{1}{(2\pi\sigma^2)^{1/2}} \exp \left\{ -\frac{1}{2\sigma^2}(x - \mu)^2 \right\}$$

What is the *likelihood* of this set of measurements ?

Can we estimate μ and σ given \mathbf{x} ?

Likelihood and regression (*)

Likelihood and parameter estimation

Since the variables \mathbf{x} are i.i.d we can write the joint probability distribution, therefore the **likelihood** of the dataset, given μ and σ is:

$$\mathcal{L}(\mu, \sigma^2; \mathbf{x}) \equiv p(\mathbf{x}|\mu, \sigma^2) = \prod_{n=1}^N \mathcal{N}(x_n|\mu, \sigma^2)$$

To **estimate** μ and σ given \mathbf{x} one **maximizes** p w.r.t these parameters.
In practice often maximize $\ln(p)$ or minimize $-\ln(p)$.

$$\ln p(\mathbf{x}|\mu, \sigma^2) = -\frac{1}{2\sigma^2} \sum_{n=1}^N (x_n - \mu)^2 - \frac{N}{2} \ln \sigma^2 - \frac{N}{2} \ln(2\pi)$$

$$\begin{cases} \frac{\partial(\ln p(\mathbf{x}|\mu, \sigma^2))}{\partial \mu} = 0 \rightarrow \mu_{\text{ML}} = \frac{1}{N} \sum_{n=1}^N x_n \\ \frac{\partial(\ln p(\mathbf{x}|\mu, \sigma^2))}{\partial \sigma} = 0 \rightarrow \sigma_{\text{ML}} = \frac{1}{N} \sum_{n=1}^N (x_n - \mu_{\text{ML}})^2 \end{cases}$$

Expected values

$$\begin{aligned} \mathbb{E}[\mu_{\text{ML}}] &= \mu \\ \mathbb{E}[\sigma_{\text{ML}}^2] &= \left(\frac{N-1}{N}\right) \sigma^2 \end{aligned}$$


Likelihood and regression (*)

Curve fitting with noise

Assume target variable in training dataset is subject to Gaussian noise

$$p(t|x, \mathbf{w}, \beta) = \mathcal{N}(t|y(x, \mathbf{w}), \beta^{-1})$$

where $\beta = \frac{1}{\sigma^2}$ is a precision parameter.



Likelihood and regression (*)

Predictive probabilistic model

By maximizing the likelihood on the training dataset we obtain a probabilistic predictive model for t (instead of a single point estimate):


$$p(t|x, \mathbf{w}_{\text{ML}}, \beta_{\text{ML}}) = \mathcal{N}(t|y(x, \mathbf{w}_{\text{ML}}), \beta_{\text{ML}}^{-1})$$

where \mathbf{w}_{ML} is obtained by minimizing the sum of square error $E(\mathbf{w})$

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2$$

and β_{ML} is given by

$$\frac{1}{\beta_{\text{ML}}} = \frac{1}{N} \sum_{n=1}^N \{y(x_n, \mathbf{w}_{\text{ML}}) - t_n\}^2$$



Exercise: show this !

Likelihood and regression (*)

Hints →

$$p(\mathbf{t}|\mathbf{x}, \mathbf{w}, \beta) = \prod_{n=1}^N \mathcal{N}(t_n | y(x_n, \mathbf{w}), \beta^{-1})$$

$$\ln p(\mathbf{t}|\mathbf{x}, \mathbf{w}, \beta) = - \underbrace{\frac{\beta}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2}_{\beta E(\mathbf{w})} + \frac{N}{2} \ln \beta - \frac{N}{2} \ln(2\pi)$$

\mathbf{w}_{ML} is determined by minimizing sum of square error $E(\mathbf{w})$, then:

$$\frac{1}{\beta_{\text{ML}}} = \frac{1}{N} \sum_{n=1}^N \{y(x_n, \mathbf{w}_{\text{ML}}) - t_n\}^2$$

Likelihood and regression (*)

Bayes theorem and likelihood



$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

Posterior knowledge
on theory

Likelihood of observing
these data given a theory

Prior knowledge
on theory

$$P(\text{theory}|\text{data}) = \frac{P(\text{data}|\text{theory})P(\text{theory})}{P(\text{data})}$$

Usually just a
normalisation factor

Likelihood and regression (*)

Bayesian approach

We assume that unknown parameters \mathbf{w} follow a Gaussian **prior** of the form:

$$p(\mathbf{w}|\alpha) = \mathcal{N}(\mathbf{w}|\mathbf{0}, \alpha^{-1}\mathbf{I}) = \left(\frac{\alpha}{2\pi}\right)^{(M+1)/2} \exp\left\{-\frac{\alpha}{2}\mathbf{w}^T\mathbf{w}\right\}$$

According to **Bayes Theorem** we have the **posterior**:

$$p(\mathbf{w}|\mathbf{x}, \mathbf{t}, \alpha, \beta) \propto p(\mathbf{t}|\mathbf{x}, \mathbf{w}, \beta)p(\mathbf{w}|\alpha)$$

Taking $-\ln(p) \rightarrow$ the maximum of the posterior is given by the minimum of:

$$\frac{\beta}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2 + \boxed{\frac{\alpha}{2} \mathbf{w}^T \mathbf{w}}$$

Regularization parameter with $\lambda = \alpha/\beta$